

# EMR Cluster execution sample

- Upload this notebook to EMR Notebooks
- Set up Kernel as PySpark

```
In [1]: ## Info to run as Emr Serverless script export to python file and change EMR_SERVERLESS_EXECUTION to False
## all input data must be executed as parameters with EMR Serverless
import sys
from datetime import datetime
```

```
### EMR Serverless Execution
EMR_SERVERLESS_EXECUTION = False ## True for emr_serverless or False for Emr Cluster - Jupyter Notebook

from pyspark.sql import SparkSession
```

VBox()  
Starting Spark application

ID	YARN Application ID	Kind	State	Spark UI	Driver log	Current session?
0	application_1679779274242_0001	pyspark	idle	<a href="#">Link</a>	<a href="#">Link</a>	✓

FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...  
SparkSession available as 'spark'.  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
In [2]: spark.version
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...  
'3.3.1-amzn-0'

```
In [3]: def fnc_validate_parameters(awsExec_v=EMR_SERVERLESS_EXECUTION):
        if (len(sys.argv) != 4) and awsExec_v:
            print("Usage: spark-etl ['input folder'] ['output folder'] ['rpt_folder']")
            sys.exit(-1)

        if not(awsExec_v):
            ## Emr Cluster execution
            input_args = ['python-script.py', '../s3_data/input/', '../s3_data/output/', '../s3_data/rpt/']
            input_location = input_args[1]
            output_location = input_args[2]
            rpt_location = input_args[3]
        else:
            ## Emr Serverless Execution
            input_location = sys.argv[1]
            output_location = sys.argv[2]
            rpt_location = sys.argv[3]

        return input_location, output_location, rpt_location
```

VBox()  
FloatProgress(value=0.0, bar\_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...

```
In [ ]: (input_location, output_location, rpt_location) = fnc_validate_parameters()
```

```
In [5]: def fshape(dataframe1):
        print('Shape : ', dataframe1.count(), len(dataframe1.columns))
```

VBox()

```
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [6]: dbname = 'DBM' ## database Marketing
        tablename = 'TBP_CUSTOMER_CLV' ## Parquet table - Customer Lifetime Value
        spark_ml_table = 'TB_ML_SPARK_SDF'
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [7]: def fnc_show_db_tables():
        spark.sql(' SHOW DATABASES ').show()
        spark.sql(' SHOW TABLES ').show()

        def spark_sql_write_glue_database(db_name, table_name, parquet_output_location=output_location, temp_table=spark_ml_table):

            ## Create AWS GLUE table for Analytics - Ad hoc query for example using Athena SQL
            print(' database creation: ', db_name)
            spark.sql(f" CREATE database if not exists {db_name} ")

            print(' table name creation , ', table_name)
            spark.sql((
                f" CREATE TABLE IF NOT EXISTS {db_name}.{table_name} "
                f" USING PARQUET LOCATION '{parquet_output_location}' AS SELECT * FROM {temp_table}"
            ))
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [8]: # s3_bucket = 's3://S3_BUCKET_NAME/'
        # parquet_file_name = 'S3_NAME_DIR/data_input/OnlineRetail__AWS.parquet'

        # s3_filename = s3_bucket + parquet_file_name

        s3_filename = input_location
        s3_filename
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
's3://s3-bucket-ohio-v1/emr_serverless/s3_data/input/'
```

```
In [9]: ## read parquet filename
        sdf = spark.read.parquet(s3_filename)
        print(sdf.printSchema())
        # fshape(sdf)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
root
 |-- InvoiceNo: string (nullable = true)
 |-- StockCode: string (nullable = true)
 |-- Description: string (nullable = true)
 |-- Quantity: integer (nullable = true)
 |-- InvoiceDate: timestamp (nullable = true)
 |-- UnitPrice: float (nullable = true)
 |-- CustomerID: double (nullable = true)
 |-- Country: string (nullable = true)
```

None

```
In [10]: fshape(sdf)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
Shape : 541909 8
```

```
In [11]: ## ETL
sdf.createOrReplaceTempView('TB_SALES_SDF')
spark.sql('select max(TO_DATE(InvoiceDate)) as current_date_for_FRMV_CLV, current_date as not_today from TB_SALES_SDF').show()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
+-----+-----+
|current_date_for_FRMV_CLV| not_today|
+-----+-----+
|                2011-12-09|2023-03-25|
+-----+-----+
```

```
In [12]: ## formula to calculate CLV
def fnc_customer_clv_udf(monetary_value_f, frequency_f, recency_f, discount_f=0.1):
    return round ( ( (monetary_value_f / frequency_f) * (1 - ((recency_f + 1) / 365)) / (1 + discount_f) ) , 2)
```

*## Register the formula to be used by Spark-SQL*

```
from pyspark.sql.types import FloatType
```

```
spark.udf.register('fnc_customer_clv_udf', fnc_customer_clv_udf, FloatType())
```

```
print("Catalog Entry:")
```

```
[print(r) for r in spark.catalog.listFunctions() if "fnc_customer_clv_udf" in r.name]
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
Catalog Entry:
Function(name='fnc_customer_clv_udf', description=None, className='org.apache.spark.sql.UDFRegistration$$Lambda$4837/1074466461', isTemporary=True)
[None]
```

```
In [13]: ## Apply some filters and create the main customer purchase history as an example
```

```
sql_query_clv = """
```

```
WITH TB_SALES_V AS
```

```
(
```

```
    SELECT CustomerID as customer_id
           , COUNT(DISTINCT (InvoiceDate)) as frequency
           , DATEDIFF( current_date , MAX (InvoiceDate) ) as recency_now
           , ROUND(SUM(Quantity * UnitPrice), 2) as monetary_value
           , ROUND(avg(Quantity * UnitPrice), 2) as avg_revenue
           , MIN(InvoiceDate) as dt_first_Invoice
           , MAX(InvoiceDate) as dt_last_Invoice
           -- , ROUND(AVG(Quantity), 2) as avg_items
           -- , ROUND(SUM(Quantity), 2) as total_items
```

```
FROM TB_SALES_SDF
```

```
WHERE 1 = 1
```

```
    AND InvoiceDate IS NOT NULL
```

```
    AND Quantity > 0
```

```
    AND UnitPrice > 0
```

```
GROUP BY customer_id
```

```
)
```

```
SELECT tb3.*
```

```
    , ROUND ( ( (monetary_value / frequency) * (1 - ((recency_dt + 1) / 365)) / (1 + 0.1) ) , 2) AS CLV_SQL -- discount of 0.1
```

```
    , fnc_customer_clv_udf(monetary_value,frequency,recency_dt) AS CLV_UDF
```

```
FROM (
```

```
    SELECT tb1.*
```

```
           , CAST( DATEDIFF(tb2.dt_current_date , tb1.dt_last_Invoice ) as float) as recency_dt
```

```
    FROM TB_SALES_V as tb1
```

```
    CROSS JOIN (SELECT MAX(dt_last_Invoice) AS dt_current_date FROM TB_SALES_V) tb2
```

```
    ) tb3
```

```
WHERE 1 = 1
```

```
    AND monetary_value > 0
```

```
    AND frequency > 0
```

```
    AND customer_id IS NOT NULL
```

```
ORDER BY monetary_value DESC
```

```
sdf_clv = spark.sql(sql_query_clv)
sdf_clv.printSchema()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
root
|-- customer_id: double (nullable = true)
|-- frequency: long (nullable = false)
|-- recency_now: integer (nullable = true)
|-- monetary_value: double (nullable = true)
|-- avg_revenue: double (nullable = true)
|-- dt_first_Invoice: timestamp (nullable = true)
|-- dt_last_Invoice: timestamp (nullable = true)
|-- recency_dt: float (nullable = true)
|-- CLV_SQL: double (nullable = true)
|-- CLV_UDF: float (nullable = true)
```

```
In [14]: print('clv_SQL and clv_udf provide the same information - just show how to implement it using 2 solutions... SQL and UDF')
         sdf_clv.show(3)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
clv_SQL and clv_udf provide the same information - just show how to implement it using 2 solutions... SQL and UDF
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|customer_id|frequency|recency_now|monetary_value|avg_revenue|dt_first_Invoice|dt_last_Invoice|recency_dt|CLV_SQL|CLV_UDF|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|14646.0|51|4125|200541.0|137.36|2010-12-20 10:09:00|2011-12-08 00:12:00|1.0|3555.12|3555.12|
|16446.0|2|4124|168472.49|56157.5|2011-05-18 09:52:00|2011-12-09 09:15:00|0.0|76368.6|76368.6|
|17450.0|27|4134|121321.71|588.94|2010-12-07 09:23:00|2011-11-29 09:56:00|10.0|3961.8|3961.8|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 3 rows
```

```
In [15]: sdf_clv.createOrReplaceTempView(spark_ml_table)
```

```
VBox()  
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [16]: def ml_sql_prediction():
text_sql_ml2 = f"""
SELECT
    {spark_ml_table}.*,
    ( CASE
WHEN ( ( `frequency` > 1.0e1 AND `frequency` <= 1.14e2 ) ) THEN 9
WHEN ( ((abs(year(`dt_first_Invoice`) - 2.01e3) <= 10e-9) OR ( `dt_first_Invoice` IS NULL ) ) ) AND ((abs(`frequency` - 1.0e0) <= 10e-9) OR (abs(`frequency` - 2.0e0) <= 10e-9) OR (abs(`freq
WHEN ( ((abs(`frequency` - 7.0e0) <= 10e-9) OR ( `frequency` >= 8.0e0 AND `frequency` <= 1.3e1 ) ) ) THEN 3
WHEN ( ( ( `recency_dt` >= 0.0e0 AND `recency_dt` <= 4.0e0 ) ) ) THEN 10
WHEN ( ( ( `CLV_SQL` > 9.0245000000000005e2 AND `CLV_SQL` <= 7.49729e3 ) ) ) THEN 6
WHEN ( ( ( `CLV_SQL` > 3.8501999999999998e2 AND `CLV_SQL` <= 9.0245000000000005e2 ) ) AND ((abs(`frequency` - 1.0e0) <= 10e-9) OR (abs(`frequency` - 2.0e0) <= 10e-9) OR (abs(`frequency` -
WHEN ( ( ( (datediff(concat(year(`dt_first_Invoice`),'-',month(`dt_first_Invoice`),'-',day(`dt_first_Invoice`)),concat(year(`dt_first_Invoice`),'-01-01')) + 1) > 1.3e1 AND (datediff(concat(
WHEN ( ((abs(month(`dt_last_Invoice`) - 3.0e0) <= 10e-9) OR (abs(month(`dt_last_Invoice`) - 4.0e0) <= 10e-9) OR (abs(month(`dt_last_Invoice`) - 5.0e0) <= 10e-9) OR (abs(month(`dt_last_Invoic
WHEN ( ( ( `recency_dt` >= 3.0e0 AND `recency_dt` <= 2.5e1 ) OR ( `recency_dt` > 3.1e1 AND `recency_dt` <= 3.6e1 ) OR ( `recency_dt` > 3.25e2 AND `recency_dt` <= 3.74e2 ) ) AND ((abs
WHEN ( ( ( (datediff(concat(year(`dt_last_Invoice`),'-',month(`dt_last_Invoice`),'-',day(`dt_last_Invoice`)),concat(year(`dt_last_Invoice`),'-01-01')) + 1) >= 4.0e0 AND (datediff(concat(yea
ELSE 11
END ) AS kc_monetary_value
FROM {spark_ml_table}

"""
return text_sql_ml2
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [17]: sdf_ml = spark.sql(ml_sql_prediction())
```

```
sdf_ml.printSchema()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
root
|-- customer_id: double (nullable = true)
|-- frequency: long (nullable = false)
|-- recency_now: integer (nullable = true)
|-- monetary_value: double (nullable = true)
|-- avg_revenue: double (nullable = true)
|-- dt_first_Invoice: timestamp (nullable = true)
|-- dt_last_Invoice: timestamp (nullable = true)
|-- recency_dt: float (nullable = true)
|-- CLV_SQL: double (nullable = true)
|-- CLV_UDF: float (nullable = true)
|-- kc_monetary_value: integer (nullable = false)
```

```
In [18]: sdf_ml.show(3)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|customer_id|frequency|recency_now|monetary_value|avg_revenue| dt_first_Invoice| dt_last_Invoice|recency_dt|CLV_SQL|CLV_UDF|kc_monetary_value|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 14646.0| 51| 4125| 200541.0| 137.36|2010-12-20 10:09:00|2011-12-08 00:12:00| 1.0|3555.12|3555.12| 9|
| 16446.0| 2| 4124| 168472.49| 56157.5|2011-05-18 09:52:00|2011-12-09 09:15:00| 0.0|76368.6|76368.6| 10|
| 17450.0| 27| 4134| 121321.71| 588.94|2010-12-07 09:23:00|2011-11-29 09:56:00| 10.0| 3961.8| 3961.8| 9|
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 3 rows
```

```
In [19]: s3_export_file = output_location
sdf_ml.write.mode('overwrite').parquet(s3_export_file)
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```
In [20]: ## Summary report
sdf_ml.createOrReplaceTempView('TB_CLV_SDF_ML')

ml_rpt_sql = """
WITH TB_CLUSTER AS
(
    select kc_monetary_value as cluster_number
    , count(distinct customer_id) as customer_count
    , avg(clv_sql) avg_clv
    , avg(monetary_value) avg_monetary_value
    -- , count(*) as qty_records
    FROM TB_CLV_SDF_ML
    group by kc_monetary_value
)
SELECT cluster_number
--      , customer_count
      , ROUND( customer_count / (select sum(customer_count) from TB_CLUSTER ) * 100, 2) as percent_of_customers
      , ROUND( avg_clv, 2) as avg_clv
      , ROUND( avg_monetary_value, 2) as avg_monetary_value
FROM TB_CLUSTER tb1
order by avg_clv desc
"""

sdf_ml_rpt = spark.sql(ml_rpt_sql)
sdf_ml_rpt.printSchema()
```

```
VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
```

```

root
|-- cluster_number: integer (nullable = false)
|-- percent_of_customers: double (nullable = true)
|-- avg_clv: double (nullable = true)
|-- avg_monetary_value: double (nullable = true)

```

In [22]: sdf\_ml\_rpt.show()

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
+-----+-----+-----+-----+
|cluster_number|percent_of_customers|avg_clv|avg_monetary_value|
+-----+-----+-----+-----+
|          6|          1.41|1627.96|          6390.09|
|         10|          4.78| 859.47|          2221.99|
|          9|          3.88| 607.03|         15079.84|
|          2|         10.07| 543.86|          1592.38|
|          3|          5.45| 359.51|          3559.05|
|          1|          8.47|  201.7|          1180.36|
|          5|         14.02| 181.48|           579.6|
|          4|         33.19| 156.01|          519.58|
|          7|         10.07| 145.49|          342.34|
|          8|          8.66|  126.8|          529.09|
+-----+-----+-----+-----+

```

## AWS GLUE - Database and Table export

In [26]: 

```
## EMR CLUSTER
### upload table to AWS GLUE - database
spark_sql_write_glue_database(db_name=dbname, table_name=tablename,
                             parquet_output_location=output_location)
```

```

VBox()
FloatProgress(value=0.0, bar_style='info', description='Progress:', layout=Layout(height='25px', width='50%'),...
database creation: DBM
table name creation , TBP_CUSTOMER_CLV

```

In [ ]: !jupyter nbconvert --to html customer\_value\_emr\_cluster\_execution\_sample.ipynb